

SRM INSTITUTE OF SCIENCE AND TECHNOLOGY

Team Details :

PRIYANSHU THAKUR (RA2211003011756)

ADITYA GUPTA (RA2211003011821)

RITESH MAHARA (RA2211003011937)



AI-Based Kubernetes Pod Failure Prediction Model

Problem Statement:

Kubernetes clusters often face issues like **pod failures**, which lead to service downtime and degraded system performance. Detecting these failures early is crucial for maintaining system stability.

Our Approach:

For this hackathon, our team focused on **predicting Kubernetes pod failures** using a Machine Learning model. After analyzing various possible issues, we found that **pod failure prediction** — especially detecting failures like **CrashLoopBackOff** — was the most feasible and impactful to solve within the timeframe.

STEP 1: Creation of a Pods failure for the visual representation

K L L K C O D A

Twitter LinkedIn GitHub PLUS

Areas Account Creator

Kubernetes 1.32

This playground will always have the latest Kubeadm Kubernetes version a few weeks after release.

You have access to an empty Kubeadm cluster with two 2GB nodes. The controlplane node has taint removed to be able to schedule workload as well.

There are also more [K8s Playgrounds](#) like [One-Node-4GB](#) available.

This is just an empty environment, if you're looking for scenarios check [CKS](#), [CKA](#), [CKAD](#) or all [Areas](#).

START

Editor Tab 1 +

Initialising Kubernetes... done

controlplane:~\$ kubectl get nodes

NAME	STATUS	ROLES	AGE	VERSION
controlplane	Ready	control-plane	45h	v1.32.1
node01	Ready	<none>	45h	v1.32.1

controlplane:~\$ nano faulty-pod.yaml

controlplane:~\$ kubectl apply -f faulty-pod.yaml

pod/faulty-pod created

controlplane:~\$ kubectl get pods --watch

NAME	READY	STATUS	RESTARTS	AGE
faulty-pod	0/1	CrashLoopBackOff	1 (2s ago)	8s
faulty-pod	0/1	Error	2 (14s ago)	20s
faulty-pod	0/1	CrashLoopBackOff	2 (12s ago)	32s
faulty-pod	0/1	Error	3 (26s ago)	46s
faulty-pod	0/1	CrashLoopBackOff	3 (11s ago)	57s
faulty-pod	0/1	Error	4 (49s ago)	95s
faulty-pod	0/1	CrashLoopBackOff	4 (13s ago)	108s
faulty-pod	0/1	Error	5 (81s ago)	2m56s
ppfaulty-pod	0/1	CrashLoopBackOff	5 (12s ago)	3m8s

kubectl describe pod faulty-pod

faulty-pod 0/1 Error 6 (2m46s ago) 5m42s

faulty-pod 0/1 CrashLoopBackOff 6 (13s ago) 5m54s

^Ccontrolplane:~\$ kubectl describe pod faulty-pod

Name: faulty-pod

Namespace: default

Priority: 0

Service Account: default

Node: node01/172.30.2.2

Start Time: Mon, 24 Mar 2025 17:28:44 +0000

Labels: <none>

Annotations: cni.projectcalico.org/containerID: 2d7aa4727e9505587384c3e481c761f1f744be0ab3e12b5509c5a000b14be017

Status: Running

IP: 192.168.1.4

Containers:

crash-container:

Container ID: containerd://f474529f5203c93db312767ab4d5b919b6ab45d4b408163376d8085eeac5573

Image: alpine

Image ID: docker.io/library/alpine@sha256:a8560b36e8b8210634f77d9f7f9efd7ffa463e380b75e2e74aff4511df3ef88c

Port: <none>

Host Port: <none>

K L L K C O D A

PLUS

AreasAccountCreatorLogout

Kubernetes1.32

This playground will always have the latest Kubeadm Kubernetes version a few weeks after release.

You have access to an empty Kubeadm cluster with two 2GB nodes. The controlplane node has taint removed to be able to schedule workload as well.

There are also more [K8s Playgrounds](#) like [One-Node-4GB](#) available.

This is just an empty environment, if you're looking for scenarios check [CKS](#), [CKA](#), [CKAD](#) or all [Areas](#).

START

Editor Tab1 +

30 min

```
^Controlplane:~$ kubectl describe pod faulty-pod
Name:          faulty-pod
Namespace:     default
Priority:       0
Service Account: default
Node:          node01/172.30.2.2
Start Time:    Mon, 24 Mar 2025 17:28:44 +0000
Labels:        <none>
Annotations:   cni.projectcalico.org/containerID: 2d7aa4727e9505587384c3e481c761f1f744be0ab3e12b5509c5a000b14be017
               cni.projectcalico.org/podIP: 192.168.1.4/32
               cni.projectcalico.org/podIPs: 192.168.1.4/32
Status:        Running
IP:            192.168.1.4
IPs:           192.168.1.4
Containers:
  crash-container:
    Container ID:   containerd://f474529f5203c93db312767abcd5b919b6ab45d4b408163376d8085eeac5573
    Image:          alpine
    Image ID:       docker.io/library/alpine@sha256:a8560b36e8b8210634f77d9ff79efd7ffa463e380b75e2e74aff4511df3ef88c
    Port:           <none>
    Host Port:      <none>
    Command:
      sh
      -c
      exit 1
    State:          Waiting
      Reason:       CrashLoopBackOff
    Last State:     Terminated
      Reason:       Error
      Exit Code:    1
    Started:        Mon, 24 Mar 2025 17:34:25 +0000
    Finished:       Mon, 24 Mar 2025 17:34:25 +0000
    Ready:          False
    Restart Count:  6
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-nspng (ro)
Conditions:
  Type                     Status
  PodReadyToStartContainers True
  Initialized              True
  Ready                    False
  ContainersReady          False
  PodScheduled             True
```

K L L K C O D A

PLUS

AreasAccountCreatorLogout

Kubernetes1.32

This playground will always have the latest Kubeadm Kubernetes version a few weeks after release.

You have access to an empty Kubeadm cluster with two 2GB nodes. The controlplane node has taint removed to be able to schedule workload as well.

There are also more [K8s Playgrounds](#) like [One-Node-4GB](#) available.

This is just an empty environment, if you're looking for scenarios check [CKS](#), [CKA](#), [CKAD](#) or all [Areas](#).

START

Editor Tab1 +

29 min

```
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-nspng (ro)
Conditions:
  Type                     Status
  PodReadyToStartContainers True
  Initialized              True
  Ready                    False
  ContainersReady          False
  PodScheduled             True
Volumes:
  kube-api-access-nspng:
    Type:               Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:       kube-root-ca.crt
    ConfigMapOptional:    <nil>
    DownwardAPI:         true
    QoS Class:           BestEffort
    Node-Selectors:      <none>
    Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                        node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason      Age      From      Message
  ---      -
  Normal    Scheduled   7m2s     default-scheduler   Successfully assigned default/faulty-pod to node01
  Normal    Pulled      6m58s    kubelet          Successfully pulled image "alpine" in 3.508s (3.508s including waiting). Image size: 36
53068 bytes.
  Normal    Pulled      6m56s    kubelet          Successfully pulled image "alpine" in 424ms (424ms including waiting). Image size: 3653
068 bytes.
  Normal    Pulled      6m42s    kubelet          Successfully pulled image "alpine" in 482ms (482ms including waiting). Image size: 3653
068 bytes.
  Normal    Pulled      6m16s    kubelet          Successfully pulled image "alpine" in 485ms (485ms including waiting). Image size: 3653
068 bytes.
  Normal    Pulled      5m27s    kubelet          Successfully pulled image "alpine" in 368ms (368ms including waiting). Image size: 3653
068 bytes.
  Normal    Created     4m6s (x6 over 6m58s)  kubelet          Created container: crash-container
  Normal    Started     4m6s (x6 over 6m58s)  kubelet          Started container crash-container
  Normal    Pulled      4m6s     kubelet          Successfully pulled image "alpine" in 457ms (457ms including waiting). Image size: 367
068 bytes.
  Warning   BackOff     107s (x26 over 6m56s)  kubelet          Back-off restarting failed container crash-container in pod faulty-pod_default(6ff8ea1d
-bd4c-4e12-9c36-298922471294)
  Normal    Pulling     82s (x7 over 7m1s)    kubelet          Pulling image "alpine"
  Normal    Pulled      81s     kubelet          Successfully pulled image "alpine" in 407ms (407ms including waiting). Image size: 3653
068 bytes.
controlplane:~$ kubectl logs faulty-pod
controlplane:~$
```

STEP 2: Dataset Creation


We generated a dataset by collecting Kubernetes pod metrics such as:


- **Timestamp**
- **Namespace**
- **Pod Status** (Running, CrashLoopBackOff, Failed, etc.)
- **Restart Count**
- **CPU Usage**
- **Memory Usage**

We simulated pod failures by:


- Increasing the restart count beyond safe thresholds
- Assigning failure statuses like **CrashLoopBackOff**
- Injecting resource-heavy conditions (high CPU and memory usage)

Feature Engineering:

We enhanced the dataset with additional features:  **Hour of the Day**

 **Day of the Week**

 **Resource Stress Score** (based on CPU > 1.5 cores or Memory > 400MB)

 **Namespace Historical Risk** (calculated from past failure data per namespace)

STEP 3: Model Training

We used a **Random Forest Classifier** to train the model.

Data was split into training and testing sets, and scaled using **StandardScaler** for better model performance.

Training Highlights:

- Balanced class weights to handle imbalanced data
- High accuracy in predicting failure-prone pods
- Evaluated using **Confusion Matrix** and **Classification Report**

Prediction Process:

For any new pod, the model:

1. Takes real-time pod metrics (CPU, Memory, Restarts, etc.)
2. Computes the resource stress and historical namespace risk
3. Predicts whether the pod is likely to fail
4. Provides the **probability of failure**

Example Output:

✓ **Will Fail:** Yes

✓ **Failure Probability:** 85.34%

```
(venv) aditya@Adityas-MacBook-Air-6 test % python -u "/Users/aditya/projects_all/test/test.py"
Classification Report:
      precision    recall  f1-score   support

     0       0.75      0.88      0.81        88
     1       0.89      0.78      0.83       112

 accuracy      0.82      0.82      0.82       200
 macro avg      0.82      0.83      0.82       200
 weighted avg      0.83      0.82      0.82       200

Confusion Matrix:
[[77 11]
 [25 87]]

Pod Failure Prediction:
Will Fail: False
Failure Probability: 38.00%
(venv) aditya@Adityas-MacBook-Air-6 test %
```

Why This Solution is Effective:

- ✓ Combines **real-time metrics** with **historical data**
- ✓ Detects potential failures **before they happen**
- ✓ Helps DevOps teams **prevent downtime**
- ✓ Scalable and can be integrated into **Kubernetes monitoring tools**

Future Scope:

- Integrate with live Kubernetes clusters for real-time predictions
- Add alerting systems based on failure probability thresholds
- Extend the model to predict other Kubernetes issues like **node failures**, **OOM kills**, etc.

